

1 Ch 1 stuff

- $[Speedup] = \frac{[Latency\ 1]}{[Latency\ 2]} = \frac{[Throughput\ 1]}{[Throughput\ 2]}$
- $[Performance\ Improvement] = [Speedup] - 1$
- $[Exec.\ Time] = [Instr.\ Count] \times [CPI] \times [Clock\ Speed]$
- $[Performance] = [Execution\ Time]^{-1}$
- $[Dynamic\ Power] \propto [Activity] \times [Capacitance] \times [Voltage]^2 \times [Frequency]$

2 ASM

- Calle saves \$s# registers, caller saves everything else
- Stack grows down

```
.data
prompt:      .ascii "Ente...rs: \n"
input:       .space 81
inputSize:   .word 80
exitMessage_p1: .ascii "Word count: "
.text
```

3 Binary Stuff

Float

”But we’re not dealing with real numbers, this is floating point baby!”

- Dec from Float: $(-1)^{[Sign\ bit]} \times (1 + [Mantissa]) \times 2^{([Exponent\ Bits] - [Bias])}$
- Float from Dec: Convert to bin sci notation, ’sub

thebias
toget

in’ to floating exponent land

- Dec to F32:
- Zero: ”00000000” exponent, all zero mantissa, sign as usual
- Inf: ”11111111” exponent, all zero mantissa, sign as usual

- NaNs: ”11111111” exponent, sign and mantissa ”left to implementer’s discretion”
- Subnorms: ”00000000” exponent, then replace the leading 1 with a zero and continue as usual. Getting increasingly smaller but less precise

Exponent chart from jan masali:

<i>binary string</i>	<i>decimal</i>	<i>exponent</i>	<i>value</i>
00000000	0	2^{-128}	$\sim 2.94 \times 10^{-39}$
00000001	1	2^{-127}	$\sim 5.88 \times 10^{-39}$
00000010	2	2^{-126}	$\sim 1.16 \times 10^{-38}$
00000011	3	2^{-125}	$\sim 2.35 \times 10^{-38}$
01111101	125	2^{-3}	0.125
01111110	126	2^{-2}	0.25
01111111	127	2^{-1}	0.5
10000000	128	2^0	1
10000001	129	2^1	2
10000010	130	2^2	4
10000011	131	2^3	8
11111100	252	2^{124}	~ 21.3 undecillion
11111101	253	2^{125}	~ 42.5 undecillion
11111110	254	2^{126}	~ 85.1 undecillion
11111111	255	2^{127}	~ 170 undecillion

Demorgans and Boolean algebra

- Xors are 1 if there’s an even number of 1s, 0 if an even number. – Thus ’Parity’
- $\neg(A \vee B) = \neg(A) \wedge \neg(B)$
- $\neg(A \wedge B) = \neg(A) \vee \neg(B)$
- You can flip the operation and flip weather they’re internally or externally negated.

4 FSM

Not that hard, just like, make sure all states and lines can be justified and are labeled.

5 Pipelining (oh god)

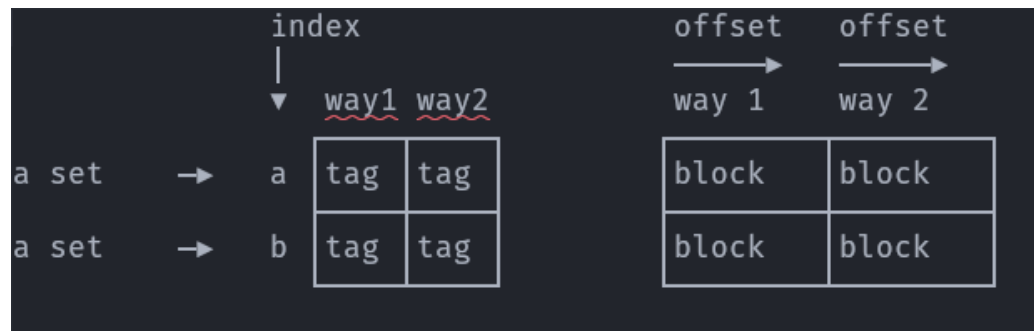
Usual steps of a 5 stage pipeline

- IF Instruction fetch 1 cycle, get next instruction from InstrMem or cache
- DR Data read from register
- AL ALU, does the computation
- DM Data memory:
- RW Read Write

Bypassing

- Point of production:
 - add, sub, etc: end of ALU
 - lw: end of DM – the mem it's reading into register
- Point of consumption:
 - add, sub, lw: start of ALU
 - sw/lw \$1, 8(\$2):
 - start of ALU for \$2
 - start of DM for \$1

6 Cache



- Offset and Index need the bits they need, tag gets the rest
 - Offset is for within the block
 - Index: which set we're referencing

- Tag array size = sets * ways * tagSize
- [Offset] = address % [block size in bytes]
- [Index] = (address / [block size in bytes]) % log₂(Sets Count)
- [Tag] = address / ([block size in bytes] * log₂([Sets Count]))
- Tag bits: remaining bits
- Offset bits: depending on blocksize, determines what byte within the block is being referenced log₂(blocksize(B)) Indexbits : determined by amount of sets log₂(setcount)

- Usually we write-allocate, bringing a miss into cache. Read misses always bring block into cache

Usually we evict the least-recently used block

- Bit string: tag index offset

7 Spectre/Meltdown

- "Spectre refers to a whole family of potential weaknesses of which meltdown is just one" - Computerphile video description
- These are the result of a combination of speculative execution and out of order execution
- Meltdown is a specific kernel memory exploit
- Need a: "lw \$t1, 0(secret); lw \$t2, \$t1" series of two instructions to leave footprints in cache

8 Virtual Memory

- Memory virtualised by the kernel, the program thinks it has a single uninterrupted area to work with
- Programs have a pagetable, which is a list of how virtual pages are mapped to physical pages in memory
- Page table translations are mostly done via the Translation Lookaside Buffer, which falls back to the pagetable

9 Multiprocessor Design

- Each core has a cache, the LLC is shared
- TODO: Protocol descriptions and examples of that chart

Cache Coherence Protocols

- Directory-based: A single location (directory) keeps track of the sharing status of a block of memory
- Snooping: Every cache block is accompanied by the sharing status of that block - all cache controllers monitor the shared bus so they can update the sharing status of the block, if necessary
- Write invalidate: a processor gains exclusive access of a block before writing by invalidating all other copies
- Write-update: when a process or writes, it updates other shared copies of that block

Locks

- Atomic exchange: simultaneous load and store operation, locks memory at the same moment it reads. – locks it while it transfers into register

MP exmaple

Request	Cache Hit/Miss	Request on the bus	Who responds	State in Cache 1	State in Cache 2	State in Cache 3	State in Cache 4
				Inv	Inv	Inv	Inv
P1: Rd X	Rd Miss	Rd X	Memory	S	Inv	Inv	Inv
P2: Rd X	Rd Miss	Rd X	Memory	S	S	Inv	Inv
P2: Wr X	Perms Miss	Upgrade X	No response. Other caches invalidate.	Inv	M	Inv	Inv
P3: Wr X	Wr Miss	Wr X	P2 responds	Inv	Inv	M	Inv
P3: Rd X	Rd Hit	-	-	Inv	Inv	M	Inv
P4: Rd X	Rd Miss	Rd X	P3 responds. Mem wrt bk	Inv	Inv	S	S

10 Basic structure of a GPU:

- many Single Instruction Multiple Thread cores, each with several warps and a warp scheduler. – large cache
- Very quick to abandon the current project if it stalls and start work on the next – minimal downtime, easy to context switch
- Each SIMT core has priavate L1 cache, large L2 shared by all cores. Each L2 bank services a subset of addresses
- The

11 Disk stuff:

- 1-12 platters (glass disks), 5-30k tracks (rings), 100-500 sectors, 512B/sector (circle around a track)
- Arm moving to correct track \Rightarrow seek time, 5-12ms, maybe less
- Rotational latency: time to rotate sector under head, usually 2ms
- Transfer time: time taken to transfer a block of bits out of disk, usually 3-65 MB/s
- Disk controller: maintains disk cache, sets up transfers.
- Mean time to Failure/Restore

Reliability MTTF

Availability: fraction of time service matches specs, $MTTF / (MTTF + MTTR)$

12 Raid types overview:

- Raid 0: No redundancy, stripes disks across drives to improve parallelism.
- Raid 1: Mirrors all disks, can sometimes increase read speed, highly redundant
- Raid 2: bit level striping, requires lockstep drives. Lots of parity drives
- Raid 3: byte level striping with dedicated parity disk. Highest sequential read speeds. Usually requires lockstep drives.
- Raid 4 and 5: block level striping, 4 has a dedicated parity disk, 5 distributes parity sections among drives.
- Raid 6: Same as raid 5, but has two parity blocks per stripe, again distributed among drives. Can work even after two drive failures. Technically raid 6 can have an arbitrary number of parity blocks added for increased redundancy.